Computational Science on Many-Core Architectures Exercise 7

Example 1 Dot Product with OpenCL (4 Points)

a)

Listing 1: opencl kernel

```
1  const char *my_opencl_program = ""
2  "__kernel void vec_mult(__global double *x,\n"
3  "                       __global double *y,\n"
4  "                       __global double *result,\n"
5  "                       unsigned int N\n)"
6  "{\n"
7  "  for (unsigned int i  = get_global_id(0);\n"
8  "                   i  < N;\n"
9  "                   i += get_global_size(0))\n"
10 "    result[i] = x[i] * y[i];\n"
11 "}";
```

This opencl kernel only multiplies the single entries and stored the results in the result vector.

Listing 2: Benchmark for the opencl

```
1  for (int i = 0;i < anz; i++)
2  {
3      //
4      // Set kernel arguments:
5      //
6      timer.reset();
7      err = clSetKernelArg(my_kernel, 0, sizeof(cl_mem),
8      (void*)&ocl_x); OPENCL_ERR_CHECK(err);
9
10     err = clSetKernelArg(my_kernel, 1, sizeof(cl_mem),
11     (void*)&ocl_y); OPENCL_ERR_CHECK(err);
12
13     err = clSetKernelArg(my_kernel, 2, sizeof(cl_mem),
14     (void*)&ocl_result); OPENCL_ERR_CHECK(err);
15
16     err = clSetKernelArg(my_kernel, 3, sizeof(cl_uint),
17     (void*)&vector_size); OPENCL_ERR_CHECK(err);
18
19     //
20     // Enqueue kernel in command queue:
21     //
22     err = clEnqueueNDRangeKernel(my_queue, my_kernel, 1,
23     NULL, &global_size, &local_size, 0, NULL, NULL); OPENCL_ERR_CHECK(err);
24
25     // wait for all operations in queue to finish:
26     err = clFinish(my_queue); OPENCL_ERR_CHECK(err);
```

```
27  }
28  opencl_time = timer.get();
29  std::cout << "Time for opencl kernel: " << opencl_time/anz << std::endl;
```

The for loop is only for the benchmark and $anz = 100$. I only add the line 13 and 14 to set the buffer for the result vector.

Listing 3: CUDA kernel

```
1  __global__ void GPU_dot (double *x, double *y, double *dot, unsigned int N)
2  {
3      unsigned int ind = threadIdx.x + blockDim.x*blockIdx.x;
4      unsigned int str = blockDim.x*gridDim.x;
5
6      double tmpsum = 0.0;
7      while(ind < N)
8      {
9          tmpsum = x[ind]*y[ind];
10         ind += str;
11     }
12     dot[threadIdx.x] = tmpsum;
13 }
```

Listing 4: CUDA benchmark

```
1  timer.reset();
2  for (int i = 0;i < anz; i++)
3  {
4      GPU_dot<<<128, 128>>>(cuda_X, cuda_Y, cuda_dot_GPU, N);
5      cudaDeviceSynchronize();
6      cudaMemcpy(dot_GPU, cuda_dot_GPU, sizeof(double), cudaMemcpyDeviceToHost);
7  }
8  //cudaMemcpy(dot_GPU, cuda_dot_GPU, sizeof(double), cudaMemcpyDeviceToHost);
9  GPU_time = timer.get();
10 std::cout << "Time for GPU kernel: " << GPU_time/anz << std::endl;
```

Listing 5: CPU kernel

```
1  void CPU_dot(ScalarType *x, ScalarType *y, ScalarType *result,
2      unsigned int N)
3  {
4      for (int i = 0; i < N; i++)
5      {
6          result[i] = x[i] * y[i];
7      }
8  }
```

Listing 6: CPU benchmark

```
1  timer.reset();
2  for (int i = 0; i < anz; i++)
3  {
4      CPU_dot(X, Y, dot_CPU, N);
5  }
6  CPU_time = timer.get();
7  std::cout << "Time for CPU kernel: " << CPU_time/anz << std::endl;
```